

— Introduction to Natural language processing (CSEG321, CSE5321, Spring 2025) —

# n-Gram Language Models

Instructor:

**Hwaran Lee**

Assistant Professor

Dept. of Artificial Intelligence

Dept. of Computer Science & Engineering

Sogang University, Seoul, South Korea



서강대학교  
SOGANG UNIVERSITY

# Class Objective

## **Understanding the concept of language models through n-gram models**

- What is an n-gram language model?
- Generating from a language model
- Evaluating a language model (perplexity)
- Smoothing (additive, interpolation, discounting)

**동해물과 백두산이** \_\_\_\_\_

**What is a n-gram language model?**

# What is a n-gram Language Model?

## Definition

- A **probabilistic model** of a sequence of words
- Joint probability distribution of words  $w_1, w_2, \dots, w_n$ :

$$P(w_1, w_2, w_3, \dots, w_n)$$

***“How likely is a given phrase, sentence, paragraph or even a document?”***

- $P(w_1, w_2, \dots, w_n)$  associated with every finite word sequence  $[w_1, w_2, \dots, w_n]$



Sample space  
= Corpus  
= Finite pieces of text

# What is a n-gram Language Model?

### Chain rule

$$\begin{aligned} & - P(w_1, w_2, w_3, \dots, w_n) \\ & = P(w_1) * P(w_2 | w_1) * P(w_3 | w_1, w_2) * P(w_4 | w_1, w_2, w_3) * \dots * P(w_n | w_1, w_2, \dots, w_{n-1}) \end{aligned}$$

conditional probability:  
 $P(w | w_1, w_2), w \in V$



- Example: "the cat sat on the mat"

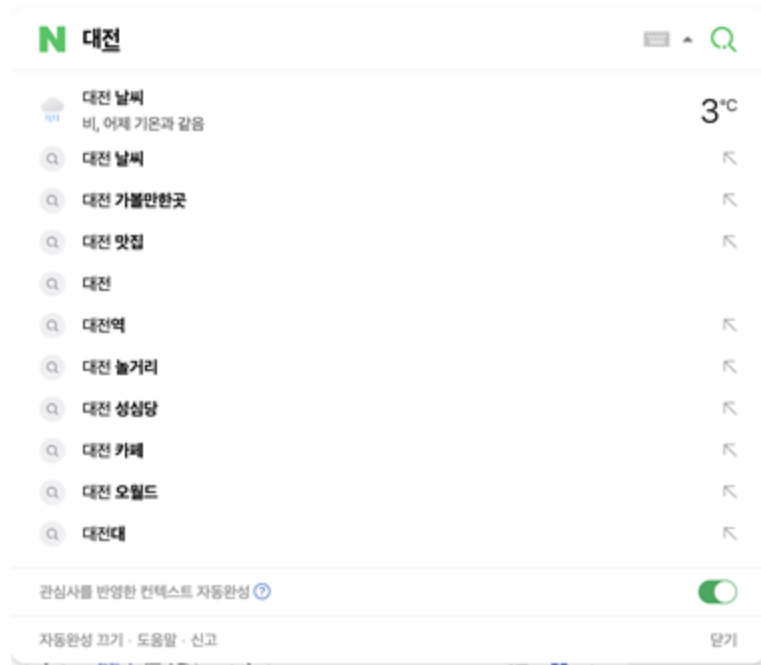
$P(\text{the cat sat on the mat})$

$$\begin{aligned} & = P(\text{the}) * P(\text{cat} | \text{the}) * P(\text{sat} | \text{the cat}) * P(\text{on} | \text{the cat sat}) * P(\text{the} | \text{the cat sat on}) \\ & * P(\text{mat} | \text{the cat sat on the}) \end{aligned}$$

# What is a n-gram Language Model?

## Chain rule

- Language models are everywhere



# What is a n-gram Language Model?

## Estimating probabilities

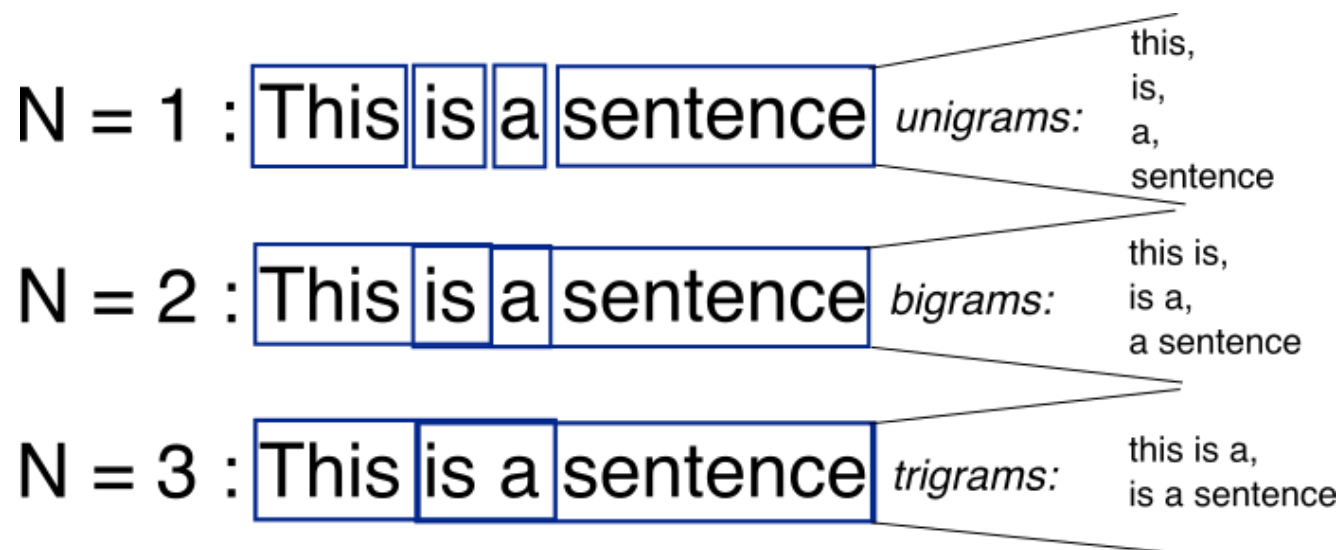
$$P(\text{sat} \mid \text{the cat}) = \text{count}(\text{the cat sat}) / \text{count}(\text{the cat})$$

bigram

$$P(\text{on} \mid \text{the cat sat}) = \text{count}(\text{the cat sat on}) / \text{count}(\text{the cat sat})$$

trigram

**Maximum  
Likelihood  
Estimation  
(MLE)!**





# What is a n-gram Language Model?

## Estimating probabilities

- Assuming we have a vocabulary of size  $V$ ,  
how many sequences of length  $n$  do we have?

a)  $n * V$

b)  $n^V$

c)  $V^n$

d)  $V/n$

# What is a n-gram Language Model?

## Estimating probabilities

- Assuming we have a vocabulary of size  $V$ ,  
how many sequences of length  $n$  do we have?  $V^n$ !
- Typical English vocabulary  $\sim 40k$  words
  - Even sentences of length  $\leq 11$  results in more than  $4 * 10^{50}$  sequences.
  - Too many to count! (# of atoms in the earth  $\sim 10^{50}$ )

# What is a n-gram Language Model?

## Markov assumption

- Use only the recent past to predict the next word
- Reduces the number of estimated parameters in exchange for modeling capacity
- 1st order
  - $P(\text{mat} \mid \text{the cat sat on the}) \approx P(\text{mat} \mid \text{the})$
- 2nd order
  - $P(\text{mat} \mid \text{the cat sat on the}) \approx P(\text{mat} \mid \text{on the})$



Andrey Markov

# What is a n-gram Language Model?

## $K^{th}$ **order Markov**

- Consider only the last k words (or less) for context

$$P(w_i | w_1, w_2, \dots, w_{i-1}) \cong P(w_i | w_{i-k}, \dots, w_{i-1})$$

which implies the probability of a sequence is:

$$P(w_1, w_2, \dots, w_n) \cong \prod_i P(w_i | w_{i-k}, \dots, w_{i-1})$$

**Need to estimate counts for up to (k+1) grams**

# What is a n-gram Language Model?

## n-gram language model

- unigram

$$P(w_1 w_2 w_3 \dots w_n) = \prod_{i=1}^n P(w_i)$$

$$\text{ex) } P(\text{the cat sat on}) = P(\text{the}) * P(\text{cat}) * P(\text{sat}) * P(\text{on})$$

- Trigram

$$P(w_1 w_2 w_3 \dots w_n) = \prod_{i=1}^n P(w_i | w_{i-1})$$

$$\text{ex) } P(\text{the cat sat on}) = P(\text{the}) * P(\text{cat} | \text{the}) * P(\text{sat} | \text{the cat}) * P(\text{on} | \text{cat sat})$$

- and 4-gram, and so on.

**Larger the n, more accurate and better the language model, but also higher costs**

**Generating from a language model**


# Generating from a Language Model

## Generation with n-gram language model

- Given a language model, how to generate a sequence?

$$\text{bigram } P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1})$$

- Generate the first word  $w_1 \sim P(w)$
- Generate the second word  $w_2 \sim P(w | w_1)$
- Generate the third word  $w_3 \sim P(w | w_2)$
- ...



# Generating from a Language Model

## Generation with n-gram language model

- Given a language model, how to generate a sequence?

$$\text{trigram } P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-2}, w_{i-1})$$

- Generate the first word  $w_1 \sim P(w)$
- Generate the second word  $w_2 \sim P(w | w_1)$
- Generate the third word  $w_3 \sim P(w | w_1, w_2)$
- Generate the fourth word  $w_4 \sim P(w | w_2, w_3)$
- ...



# Generating from a Language Model

### Generation results

- unigram

release millions See ABC accurate President of  
Donald Will cheat them a CNN megynkelly experience @  
these word out- the

- bigram

Thank you believe that @ ABC news, Mississippi  
tonight and the false editorial I think the great  
people Bill Clinton

- trigram

We are going to MAKE AMERICA GREAT AGAIN!  
#MakeAmericaGreatAgain <https://t.co/DjkdAzT3WV>

**Typical LMs are not sufficient to handle long-range dependencies**

# Generating from a Language Model

## Generation results

- Example of ChatGPT

**Do you know how to write good papers?** Writing good papers requires a blend of skills. Firstly, grasp the topic thoroughly; research extensively for credible sources. Next, organize your thoughts logically, crafting a clear thesis statement. Ensure coherence and flow in your writing, supporting arguments with evidence. Pay attention to grammar, punctuation, and style. Finally, revise and edit meticulously for clarity and effectiveness. Practice and feedback refine these skills.

prompt a.k.a conditional context

$$P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1024}, \dots, w_{i-2}, w_{i-1})$$

**Modern LMs can handle much longer contexts!**

# Generating from a Language Model

## Generation methods

- Greedy: choose the most likely word
  - To predict the next word given a context of two words  $w_1, w_2$ :

$$w_3 = \operatorname{argmax}_{w \in V} P(w | w_1, w_2)$$

```
>>> # With 'top_k' sampling, the output gets restricted the k most likely
>>> # Pro tip: In practice, LLMs use 'top_k' in the 5-50 range.
>>> outputs = model.generate(**inputs, do_sample=True, top_k=2)
```

- Top-k vs top-p sampling (also called nucleus sampling):

Token	Probability score
for	0.4
to	0.25
with	0.17
and	0.13
by	0.05


Token	Probability score
for	0.4
to	+ = 0.65 0.25
with	0.17
and	0.13
by	0.05

# **Evaluating a language model**

# Evaluating a Language Model

## Extrinsic evaluation

- Train LM apply to task observe accuracy
- Directly optimized for downstream applications
  - Higher task accuracy -> better model
- Expensive, time consuming
- Hard to optimize downstream objective (indirect feedback)



# Evaluating a Language Model

## **Intrinsic evaluation of language models**

- Research process:
  - Train parameters on a suitable training corpus
    - Assumption: observed sentences ~ good sentences
  - Test on different, unseen corpus
    - If a language model assigns a higher probability to the test set, it is better
- Evaluation metric - perplexity

## Evaluating a Language Model

### Perplexity (ppl)

- Measure of how well a LM predicts the next word
  - For a test corpus with words  $w_1, w_2, \dots, w_n$

$$\text{Perplexity} = P(w_1, w_2, \dots, w_n)^{-\frac{1}{n}}$$

$$\text{ppl}(S) = e^x \text{ where } x = -\frac{1}{n} \log P(w_1, \dots, w_n) = -\frac{1}{n} \sum_{i=1}^n \log P(w_i | w_1 \dots w_{i-1})$$

- Unigram model:  $x = -\frac{1}{n} \sum_{i=1}^n \log P(w_i)$
- Minimizing perplexity ~ maximizing probability of corpus

→ Cross-entropy

## Evaluating a Language Model

### Intuition on perplexity

- If our k-gram model (with vocabulary  $V$ ) has following probability:

$$P(w_i | w_{i-k+1}, \dots, w_{i-1}) = \frac{1}{|V|}, \quad \forall w \in V$$

- what is the perplexity of the test corpus?

A)  $e^{|V|}$       B)  $|V|$       C)  $|V|^2$       D)  $e^{-|V|}$

$$ppl(S) = e^x \text{ where } x = -\frac{1}{n} \log P(w_1, \dots, w_n) = -\frac{1}{n} \sum_{i=1}^n \log P(w_i | w_1 \dots w_{i-1})$$



Cross-entropy



## Evaluating a Language Model

### Intuition on perplexity

- If our k-gram model (with vocabulary  $V$ ) has following probability:

$$P(w_i | w_{i-k+1}, \dots, w_{i-1}) = \frac{1}{|V|}, \quad \forall w \in V$$

- what is the perplexity of the test corpus?

A)  $e^{|V|}$       B)  $|V|$       C)  $|V|^2$       D)  $e^{-|V|}$

$$\begin{aligned} ppl(S) &= e^x \text{ where } x \\ &= -\frac{1}{n} \sum_{i=1}^n \log P(w_i | w_1 \dots w_{i-1}) \end{aligned}$$

$$ppl(S) = e^{-\frac{1}{n} n \log(\frac{1}{|V|})} = |V|$$

# Evaluating a Language Model

## Perplexity

- Training corpus 38 million words, test corpus 1.5 million words, both WSJ

n-gram	unigram	bigram	trigram
perplexity	962	170	109


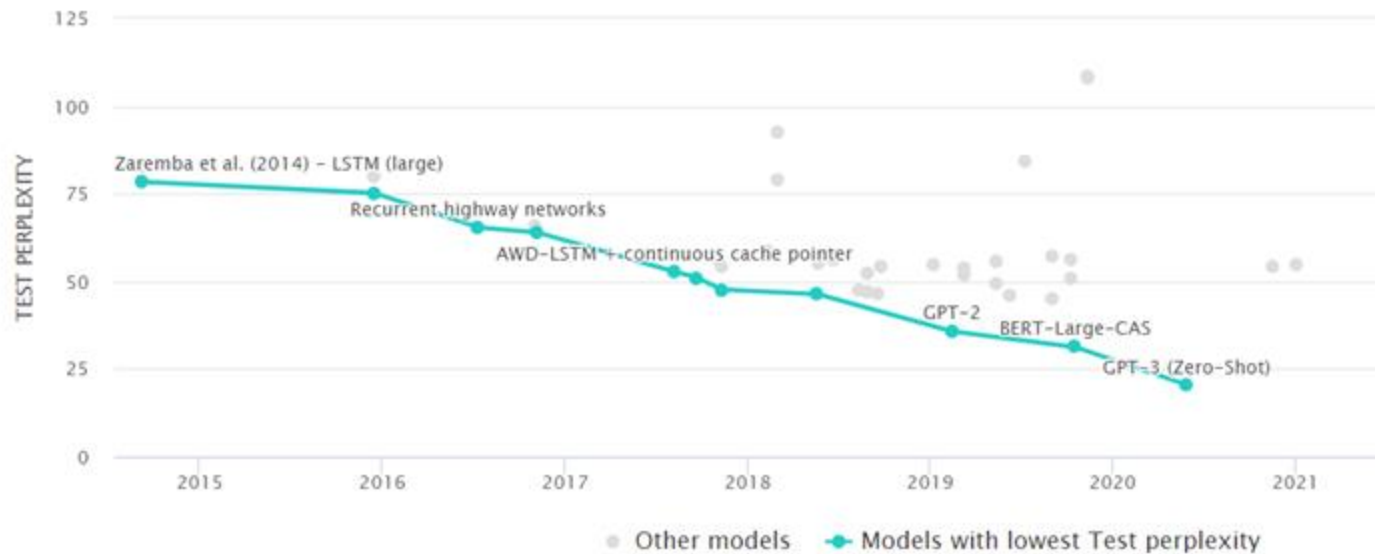


Figure 5: Validation PPL of different sized models. The larger model achieves a better performance in fewer steps, consistent with other works (Kaplan et al., 2020; Li et al., 2020).


# Smoothing

# Smoothing

## Generalization of n-grams

- Any problems with n-gram models and their evaluation?
- Not all n-grams in the test set will be observed in training data
- Test corpus might have some that have zero probability
  - Training set: Google news
  - Test set: Shakespeare
- $P(\text{affray} \mid \text{voice doth us}) = 0 \rightarrow P(\text{test corpus}) = 0$
- Perplexity is not defined

## Smoothing


**Sparsity in language**

frequency  $\propto \frac{1}{rank}$   
Zipf's Law

- Long tail of infrequent words
- Most finite-size corpora will have this problem.

# Smoothing

## Sparsity in language



frequency  $\propto \frac{1}{rank}$   
Zipf's Law

- Long tail of infrequent words
- Most finite-size corpora will have this problem.

# Smoothing

## Concept of Smoothing

- Handle sparsity by making sure all probabilities are non-zero in our model
  - **Additive**: Add a small amount to all probabilities
  - **Interpolation**: Use a combination of different granularities of n-grams
  - **Discounting**: Redistribute probability mass from observed n-grams to unobserved ones

## Smoothing

## Smoothing intuition

- When we have sparse statistics:

$P(w \mid \text{denied the})$


3 allegations

2 reports

1 claims

1 request

7 total



- Steal probability mass to generalize better

$P(w \mid \text{denied the})$

2.5 allegations


1.5 reports

0.5 claims

0.5 request

2 other

7 total





# Smoothing

## Laplace smoothing

- Also known as add-alpha
- Simplest form of smoothing: Just add  $\alpha$  to all counts and renormalize!
- Max likelihood estimate for bigrams:

$$P(w_i|w_{i-1}) = \frac{C(w_{i-1}, w_i)}{C(w_{i-1})}$$

- After smoothing:

$$P(w_i|w_{i-1}) = \frac{C(w_{i-1}, w_i) + \alpha}{C(w_{i-1}) + \alpha|V|}$$

## Smoothing

### Ray bigram counts (Berkeley restaurant corpus)

- Out of 9,222 sentences

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

## Smoothing

### Smoothed bigram counts (Berkeley restaurant corpus)

- Out of 9,222 sentences

	i	want	to	eat	chinese	food	lunch	spend
i	6	828	1	10	1	1	1	3
want	3	1	609	2	7	7	6	2
to	3	1	5	687	3	1	7	212
eat	1	1	3	1	17	3	43	1
chinese	2	1	1	1	1	83	2	1
food	16	1	16	1	2	5	1	1
lunch	3	1	1	1	1	2	1	1
spend	2	1	2	1	1	1	1	1

## Smoothing

## Smoothed bigram probabilities (Berkeley restaurant corpus)

$$P(w_i|w_{i-1}) = \frac{C(w_{i-1}, w_i) + \alpha}{C(w_{i-1}) + \alpha|V|}, \quad \alpha = 1$$

	i	want	to	eat	chinese	food	lunch	spend
i	0.0015	0.21	0.00025	0.0025	0.00025	0.00025	0.00025	0.00075
want	0.0013	0.00042	0.26	0.00084	0.0029	0.0029	0.0025	0.00084
to	0.00078	0.00026	0.0013	0.18	0.00078	0.00026	0.0018	0.055
eat	0.00046	0.00046	0.0014	0.00046	0.0078	0.0014	0.02	0.00046
chinese	0.0012	0.00062	0.00062	0.00062	0.00062	0.052	0.0012	0.00062
food	0.0063	0.00039	0.0063	0.00039	0.00079	0.002	0.00039	0.00039
lunch	0.0017	0.00056	0.00056	0.00056	0.00056	0.0011	0.00056	0.00056
spend	0.0012	0.00058	0.0012	0.00058	0.00058	0.00058	0.00058	0.00058

# Smoothing

## Linear interpolation

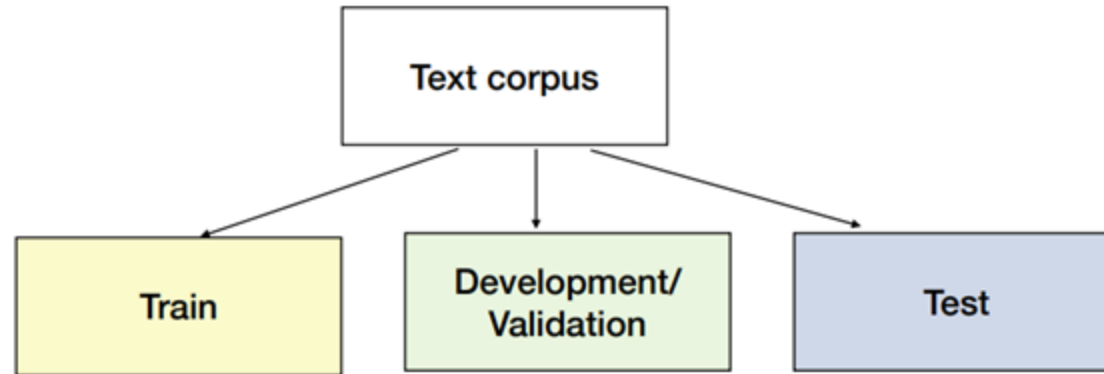
$$\hat{P}(w_i | w_{i-2}, w_{i-1}) = \lambda_1 P(w_i | w_{i-2}, w_{i-1}) + \lambda_2 P(w_i | w_{i-1}) + \lambda_3 P(w_i)$$

$$\sum_i \lambda_i = 1$$

- Use a combination of models to estimate probability
- Strong empirical performance

## Smoothing

### Linear interpolation – How can we choose lambdas?



- First, estimate n-gram probabilities on training set
- Then, estimate lambdas (as hyperparameters) to maximize probability on the held-out development/validation set
- Use best model from above to evaluate on test set

# Smoothing

## Discounting

- Determine some “mass” to remove from probability estimates
- More explicit method for redistributing mass among unseen n-grams
- Just choose an absolute value to discount (usually  $< 1$ )

# Smoothing

## Absolute Discounting

- Define  $\text{Count}^*(x) = \text{Count}(x) - 0.5$

- Missing probability mass:

$$\alpha(w_{i-1}) = 1 - \sum_w \frac{\text{Count}^*(w_{i-1}, w)}{\text{Count}(w_{i-1})}$$

$$\alpha(\text{the}) = 10 \times \frac{0.5}{48} = \frac{5}{48}$$

- Divide this mass between words  $w$  for which  $\text{Count}(\text{the}, w) = 0$

$x$	$\text{Count}(x)$	$\text{Count}^*(x)$	$\frac{\text{Count}^*(x)}{\text{Count}(x)}$
the	48		
the, dog	15	14.5	14.5/48
the, woman	11	10.5	10.5/48
the, man	10	9.5	9.5/48
the, park	5	4.5	4.5/48
the, job	2	1.5	1.5/48
the, telescope	1	0.5	0.5/48
the, manual	1	0.5	0.5/48
the, afternoon	1	0.5	0.5/48
the, country	1	0.5	0.5/48
the, street	1	0.5	0.5/48



## Smoothing

**Absolute Discounting**

$$\alpha(the) = 10 \times \frac{0.5}{48} = \frac{5}{48}$$

$$P_{abs-discount}(w_i | w_{i-1})$$

$$= \begin{cases} \frac{c(w_{i-1}, w_i) - d}{c(w_{i-1})}, & \text{if } c(w_{i-1}, w_i) > 0 \\ \alpha(w_{i-1}) \cdot \frac{P(w_i)}{\sum_{w'} P(w')}, & \text{if } c(w_{i-1}, w_i) = 0 \end{cases}$$

$x$	$\text{Count}(x)$	$\text{Count}^*(x)$	$\frac{\text{Count}^*(x)}{\text{Count}(x)}$
the	48		
the, dog	15	14.5	14.5/48
the, woman	11	10.5	10.5/48
the, man	10	9.5	9.5/48
the, park	5	4.5	4.5/48
the, job	2	1.5	1.5/48
the, telescope	1	0.5	0.5/48
the, manual	1	0.5	0.5/48
the, afternoon	1	0.5	0.5/48
the, country	1	0.5	0.5/48
the, street	1	0.5	0.5/48

***E.O.D***